
Cortical Decoders Trained on Other Individuals Age Slower Than Your Own

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 With adapting brain–computer interface (BCI) technology, neural decoding algo-
2 rithms are one of the most essential parts of being able to understand people’s
3 brains. Decoders must work for months, and recalibrating them per subject has high
4 subject costs. A decoder fit to a singular animal’s cortex drifts, losing accuracy over
5 days, and usually the remedy is recalibrating on fresh data from that same subject.
6 We report a different kind of robustness: a decoder built from one other individual,
7 or pooled from a set of other individuals, ages more slowly than the subject’s
8 own. On the open BraiDyn-BC cohort (25 mice, 44 Allen-atlas cortical parcels, a
9 ~15-day operant protocol) we decode four task events and measure, per mouse,
10 the AUC lost over two weeks by its own early-week decoder against one pooled
11 over the other 24 mice and one built from a single other mouse. The personal
12 decoder ages more by +0.017 AUC ($p = 5 \times 10^{-5}$, 73 of 100 mouse–event pairs).
13 This also is not a data-volume effect: at a matched training-event count a decoder
14 from a single other mouse still ages less. Hence we conclude this is rather an
15 effect of not overfitting an individual’s drift-prone idiosyncrasies. It also survives
16 a 1-D CNN and a GRU. The effect holds past this cohort and recording modality.
17 On an independent two-photon dataset from a different laboratory – Allen Visual
18 Behavior, 23 mice, single-plane VISp populations decoding a stimulus change –
19 the same session-held-out protocol gives +0.023 AUC ($p = 0.003$, 74% of mice),
20 and the count-matched control reproduces there as well (self – one other mouse
21 = +0.011 AUC). The early-week accuracy cost of decoding across individuals is
22 repaid as late-week stability, which suggests population priors as an alternative to
23 per-subject recalibration. A second offering is the session-held-out protocol under
24 which all of this is measured: because only the within-subject arm can share a
25 recording session with its test data, a trial-level split inflates that arm alone, and
26 holding out whole sessions is what allows this form of robustness to be reported
27 reliably. Code and a streamed-feature pipeline are released.

28 1 Introduction

29 Neural decoding is one of the most essential ways to understand neural interfaces and create BCI
30 technology, and a decoder that is deployed is a decoder that has to keep working: over months, on
31 the same individual, without fresh labelled data arriving to prop it up. Representational drift stands
32 directly against that. The reorganization of an animal’s neural code over days with fixed behavior is
33 now documented across sensory, motor, and association cortex and in the hippocampus [18, 3, 20].
34 This has a stark impact on neural decoding: a decoder trained on an animal today degrades on that
35 same animal weeks later. The dominant response is recalibration, periodically refitting the decoder on
36 fresh within-subject data [2, 5], which requires ongoing labelled data from the individual and treats
37 each subject in isolation.

The code’s cross-individual structure offers an alternate route to temporal robustness, and the field already relies on it. Cortex-wide behavioral representations are considerably shared among individuals [11, 7, 8]; latent dynamics are preserved across animals well enough that a decoder fit to one transfers to another [17]; and pooled decoders already read out behavior from unseen individuals [21, 12, 13]. Methods such as Meta-AlignNN [10] are built on that premise. While people currently know that conservation across individuals is established, and that pooling is known to work at a fixed time, every one of these results is measured at a fixed time. We offer pooling as a solution to a new problem: the fact that it resists drift over time.

Our core discovery is that a decoder that never saw the subject cannot fit to that subject’s drift-prone idiosyncrasies, so it should age more slowly, and the way we measure that difference is as an asymmetry between the two arms. This prediction, to our knowledge, has not been examined.

We test it on BraiDyn-BC [6], a widefield calcium-imaging resource whose ~ 15 -day operant protocol and shared 44-region Allen parcellation make within-subject and cross-subject decoders directly comparable across time. Our contributions:

- **Pooling resists drift.** Across a two-week period the population-pooled decoder loses less accuracy than the individual’s own, with a paired asymmetry of $+0.017$ AUC ($p = 5 \times 10^{-5}$, 73 of 100 pairs; Section 4.1). The size of this effect varies among the four events, but those are chained within a trial and have highly differing counts, so we treat the pooled estimate as the result and leave the per-event ranking uninterpreted (Section 4.2).
- **It is not a data-volume artifact.** At a matched training-event count a decoder trained on one other mouse ages less than one trained on the subject’s own, by $+0.010$ AUC on average. The gain therefore comes from training on individuals other than the subject, not from the larger training set pooling usually brings, and training on many other mice rather than one lowers aging further on all four events (Sections 4.2–4.2).
- **It replicates on an independent dataset.** On the Allen Visual Behavior two-photon cohort the same protocol gives $+0.023$ AUC ($p = 0.003$, 74% of 23 mice), with the count-matched ordering also reproduced (Section 4.4).
- **Session-level validation is load-bearing.** As trials within a recording session are dependent, a trial-level split inflates the personal decoder but not the pooled one, one-sidedly inflating the asymmetry. We report that inflation for each event; on lever pull it accounts for the entire apparent effect. For longitudinal decoding comparisons we recommend session-held-out validation as a default (Section 4.3).

2 Related work

Representational drift and its mitigation. Drift is characterized within individual subjects [18, 3, 20, 16], and most approaches to mitigating it also operate within the subject. These include manifold stabilization [2], supervised and self-calibrating recalibration [5, 15], and adversarial or nonlinear alignment to the subject’s own reference day [4, 14]. Each method uses later data from the same subject to stabilize the decoder. We instead ask whether data from other subjects can provide that stability.

Cross-subject decoding. Shared structure across individuals already supports cross-subject readout through manifold alignment [9], contrastive embeddings [19], and, on this cohort, atlas-aligned foundation models [21]. Multi-participant pooling can also reduce overfitting: HTNet, for example, pools participant data to avoid fitting a decoder to a small single-subject sample [13], and pooled pretraining across 83 animals followed by fine-tuning on held-out animals reaches state-of-the-art encoding and decoding on a standardized cross-subject cohort [12]. Our question differs in two ways. Primarily, prior work measures accuracy or transfer at a fixed time, while we look at how that accuracy changes over days, to see whether pooled decoders age more slowly. Further, we keep the training-set size fixed with the count-matched control, hence being able to delegate the relevant overfitting not to limited data but to idiosyncratic features that cause drift.

The closest work to ours combines both axes: Meta-AlignNN [10] uses the consistency of neural population dynamics across individuals to build a decoder that remains stable across subjects, time and tasks, and demonstrates it over two years in three monkeys. That work and ours point in the same practical direction, and differ in kind rather than in topic. Meta-AlignNN proposes a method for

achieving stability; we measure whether a difference in aging rate exists at all between a decoder trained on the subject and one trained on others, with the training-set size held fixed so that the comparison isolates whose data the decoder saw. A measured effect of this kind is what would justify the design choice such methods already make.

Conservation of cortex-wide dynamics. Cortex-wide activity contains shared, low-dimensional, movement-related components [11, 7, 8], and latent dynamics are preserved across animals performing similar behaviour to the point that a decoder fit to one animal transfers to another [17]. That conservation is the mechanism a pooled decoder would exploit, and it is why pooling is possible at all here. It does not, however, imply that pooling will *resist drift*: a code can be shared across individuals and still move within each of them at the same rate. Whether the shared component is also the more temporally stable one is the empirical question we ask.

3 Data and methods

3.1 Datasets and preprocessing

BraiDyn-BC [6] (DANDI:001425, CC-BY) provides, for each mouse and session, hemodynamically corrected dorsal-cortex $\Delta F/F$ traces parcellated into 44 Allen Common Coordinate Framework regions at ~ 30 Hz, together with time-aligned behavioral channels. The dataset includes 25 mice, each of which completed ~ 15 daily operant sessions over ~ 3 weeks. We stream the 44-region traces and behavioral channels without downloading the raw movies.

The two cohorts used in this paper were built for different purposes, and the ways they differ are what make the comparison informative rather than redundant. BraiDyn-BC is densely longitudinal, with roughly fifteen sessions per mouse over three weeks, and it supplies a shared coordinate system: the 44 Allen parcels mean that region k in one mouse is the same region in every other, so a pooled decoder can be fit directly in a common feature space. The Allen Visual Behavior cohort has neither property. Its longitudinal subset is sparser, and because it is cellular two-photon imaging there is no correspondence between animals at all – one mouse’s cell 17 is unrelated to another’s – so a common feature space has to be constructed rather than assumed (Section 3.4). The recording modality, the task, and the decoded contrast all differ as well. A result that survives both is therefore not a property of one atlas, one instrument, or one behavior.

3.2 Features and decoders

We detect the onset of four events (lever pull, lick, reward, and tone) as rising edges, with a 1 s refractory period. For each onset we construct a 44-dimensional evoked feature by subtracting the mean pre-onset baseline (-0.5 – 0 s) from the mean post-onset $\Delta F/F$ response (0 – 1 s) in each parcel. Negative samples are drawn from matched-count windows more than 2 s from any onset *of the same event*; they are not screened against the other three events, and Section 4.2 reports how often they contain one. Every classification is a binary decode of one event against these matched negatives, scored by area under the ROC curve (AUC). The four decoders are therefore four independent binary detectors, not a four-way classifier: nothing in this paper asks a model to distinguish one event from another.

The primary decoder is a standardized ℓ_2 -regularized logistic regression on the 44-dimensional evoked feature. This is a linear readout of the parcel-space code: it recovers the component of behavior that is linearly present in the population average, and it is the natural comparison to prior cross-subject decoders, which are overwhelmingly linear [9, 13]. Every condition below uses this same readout, so the only variable is which data the decoder was trained on.¹

¹The result is not a property of the linear readout. Repeating the analysis with a 1-D convolutional network over time and a GRU over the frame sequence (full evoked window, -0.5 to $+1.0$ s, 45×44), under the same session-held-out protocol, the asymmetry survives for both events that carry it: lick gives $+0.020$, $+0.012$ and $+0.010$ and reward $+0.029$, $+0.015$ and $+0.012$ under the linear, CNN and GRU readouts, all significant. Under a trial-level split the nonlinear estimates appear *larger* instead, which is the inflation of Section 4.3: a higher-capacity model exploits it more effectively than a linear one.

3.3 Temporal blocks, aging, and the session-held-out protocol

We group each mouse’s sessions into an early block (days 1–5) and a late block (days 11–15) and report AUC under four conditions. The personal decoder is estimated by leaving out one whole early session at a time: for each early session s we train on the mouse’s remaining early sessions and score the resulting decoder both on s (within-mouse same-block, WS) and on the late block (within-mouse across-block, WX), then average over s . Because WS and WX come from the same fitted decoders, they differ only in which data they are tested on. The pooled decoder is trained once on the other mice’s early blocks and scored on the same held-out sessions s (LS) and on the late block (LX).

Holding out sessions rather than trials is essential and not a detail. Trials within one recording share illumination and hemodynamic state, slow arousal drift, and temporal autocorrelation, so a trial-level split lets the personal decoder train and test inside the same session. The pooled decoder is trained on other animals entirely and cannot benefit this way, so the inflation is one-sided and lands directly on the quantity of interest; Section 4.3 reports what it costs per event.

Aging is the accuracy lost when testing moves from the held-out mouse’s early data to its late block while training stays fixed: personal aging is $WS - WX$ and pooled aging is $LS - LX$. The per-mouse aging asymmetry is $(WS - WX) - (LS - LX)$; a positive value means the personal decoder ages more. We assess significance with a paired sign-flip permutation test across mice (20,000 permutations) and compute 95% confidence intervals with a cluster bootstrap over mice.

3.4 The external validation cohort

Section 4.4 repeats the analysis on the Allen Visual Behavior two-photon dataset [1]. We select containers spanning at least 14 days whose experiments are all in VISp within a single 50 μm depth bin, which is the cross-subject correspondence available here; the largest such band, 150–199 μm , holds 25 mice and 184 experiments. We require at least 15 segmented cells and 40 usable trials per experiment, leaving 166 experiments and 23 mice with at least four sessions. Sessions are streamed from S3 rather than downloaded.

As different animals have no neuron correspondence and the longitudinal subset is single-plane, averaging over the field of view would give a one-dimensional feature. The representation that is comparable across animals is therefore a permutation-invariant summary of the population response: per neuron we take the mean $\Delta F/F$ over 0–1 s minus the mean over -0.5 –0 s, then describe the distribution of those evoked responses by 25 quantiles plus its mean and standard deviation, giving a fixed 27-dimensional feature regardless of how many cells were segmented. The decoded contrast is is_change , a real image change against a sham catch trial, which is a stimulus contrast rather than a reward- or lick-coupled one. Each mouse’s sessions are split by date into early and late halves, and the personal decoder, the pooled decoder, aging, the asymmetry and the count-matched control are all defined exactly as above.

4 Results

4.1 The effect: a pooled decoder ages more slowly, by two different estimators

For each mouse we keep the test data fixed (first the mouse’s held-out early sessions, then its late block) and change only the source of the training data: either the mouse itself or the other 24 mice. Across all 100 mouse–event pairs the asymmetry is $+0.017$ AUC ($p = 5 \times 10^{-5}$), with the personal decoder aging more in 73% of pairs (Table 1).

The effect differs among the four events: it is largest for reward ($+0.029$, $p = 0.0005$) and lick ($+0.020$, $p = 0.002$), marginal for tone ($+0.015$, $p = 0.08$), and absent for lever pull ($+0.003$, $p = 0.37$). While we report this variation, we do not interpret it, as the four events all fall within roughly the same two seconds of a trial. Accordingly, the pooled estimate over all 100 pairs is the claim we defend, and the per-event column should be read as data rather than as four separate findings.

The block split compares two windows, but the effect should also appear as a continuous trend across all fifteen days. For each session day we compute the accuracy of the personal decoder and the pooled decoder and regress each on day; on early days the personal decoder is trained by

Table 1: Personal vs. pooled aging across days (AUC lost from early to late test, holding training fixed), per event, $n = 25$ mice, with whole early sessions held out. Asymmetry = personal – pooled aging; a positive value means the mouse’s own decoder ages more. p : paired sign-flip permutation. The rightmost column gives the same asymmetry under a trial-level split, the protocol this analysis replaces (Section 4.3).

Event	Personal	Pooled	Asymmetry (95% CI)	p	own ages more	trial-split
Lever-pull	−0.022	−0.025	+0.003 [−0.004, 0.011]	0.37	64%	+0.014
Lick	+0.012	−0.009	+0.020 [0.008, 0.032]	0.0023	80%	+0.028
Reward	+0.017	−0.012	+0.029 [0.014, 0.046]	0.0005	80%	+0.032
Tone	+0.016	+0.001	+0.015 [0.001, 0.032]	0.077	68%	+0.022
Pooled			+0.017 [0.010, 0.024]	5×10^{-5}	73%	+0.024

Decoding across the 15-day protocol: the personal decoder starts ahead, the pooled decoder gains faster, so the gap narrows




Figure 1: Decoding accuracy across the 15-day protocol for the personal decoder (own early data) and the pooled decoder (other mice’s early data), with per-day points and linear fits. The personal decoder starts ahead, and the pooled decoder gains more over days, so the gap between them narrows. The narrowing is significant for lick and reward, the two events where the personal decoder also loses accuracy over days, and is not significant for lever pull or tone. The per-event gap slope and its p (paired sign-flip over mice) are shown.

184 leave-one-session-out, so it is never tested on a session it trained on. The personal decoder starts
185 ahead, leading by about 0.035 AUC on day 1 as expected for a decoder trained on the same animal.
186 That lead then erodes: the gap narrows by 0.0014 AUC per day ($p = 5 \times 10^{-5}$; narrowing in 65% of
187 pairs), so by day 15 roughly 60% of the early-week advantage is gone.

188 Averaged over the four events, the pooled decoder gains accuracy over days (+0.0015 AUC/day)
189 as the sessions become easier – the mouse is more practiced – while the personal decoder stays flat
190 (+0.0000 AUC/day). The personal decoder, fit to idiosyncrasies, is unable to ride that wave, and
191 the lift it fails to take is what opens the gap. Per event, the personal decoder does lose accuracy for
192 lick (−0.0010 AUC/day) and reward (−0.0016 AUC/day), and these are the two events where the
193 narrowing is individually significant ($p = 0.007$ and $p = 0.001$; Fig. 1).

194 4.2 Ruling out the alternative explanations

195 The four decoded events come from a single cued lever-pull trial structure, so they are chained in
196 time rather than independent. Every reward onset has a lever-pull onset a median of 0.083 s away
197 and a lick onset 0.117 s away, and 100% and 98% respectively fall inside the window from which a
198 feature vector is built. Across all four events, the fraction of positive windows containing at least
199 one other event is 100% (reward), 84% (lick), 84% (tone) and 68% (lever pull). Consequently, the
200 reward decoder is never reading reward in isolation: its input window always contains licking and
201 lever movement as well. The negatives are not event-free either: screened only against the event’s
202 own onsets, they contain another event 33–57% of the time.




Figure 2: **Left:** count-matched control. With equal numbers of training events, decoder aging (AUC lost from early to late; lower values indicate greater robustness) is compared for the subject’s own data (self), one other mouse, and many other mice. **Right:** pooled aging as a function of the number of training mice N . Greater diversity provides a small additional gain, with immediate saturation for lever pull.

203 The events also differ heavily in sample size, from a median of 56 reward onsets per session against
 204 107 (tone), 530 (lick) and 652 (lever pull). The per-event asymmetry runs roughly opposite to that
 205 count; a decoder fit to fewer events has more idiosyncrasy available to overfit, and hence more
 206 for pooling to avoid. With only four events, however, we cannot separate that explanation from a
 207 behavioral one, and we therefore do not attempt to. This is why the pooled estimate over all 100 pairs
 208 is what we defend, and why later sections that select reward and lick are consistency checks on the
 209 same events rather than evidence from independent behaviors.

210 A pooled decoder contains $\sim 24\times$ more training events than a personal decoder, so its stability could
 211 reflect data volume rather than pooling. We test this with a count-matched control (Fig. 2, left). For
 212 each held-out early session s of each mouse we set the training-set size to n , the number of events in
 213 that mouse’s remaining early sessions, and compare three decoders: *self*, trained on those n events;
 214 *one*, trained on n events from a single other mouse; and *many*, trained on n events drawn across the
 215 other 24 mice. All three are scored on the same held-out session s and the same late block, so they
 216 differ only in whose data they were fit on.

217 For the events where there was an asymmetry, the count-matched ordering reproduces it. The self
 218 decoder ages more than a decoder trained on one other mouse by +0.010 AUC on average, and by
 219 +0.007 for lick, +0.017 for reward and +0.006 for tone individually. As both decoders see the same
 220 number of training events, this is not a data-volume artifact.

221 The diversity ordering is the most robust result in this analysis: *many* ages less than *one* for all four
 222 events, including lever pull (−0.026 vs. −0.020). Training on a diverse set of individuals rather than
 223 a single other individual therefore provides a smaller, additional benefit that does not depend on
 224 which event is decoded.

225 We next vary the number of training mice, $N \in \{1, 2, 4, 8, 16, 24\}$ (Fig. 2, right). As N increases,
 226 pooled aging becomes slightly more favorable for lick (−0.001 to −0.010), reward (−0.004 to
 227 −0.010), and tone (+0.003 to −0.010); lever pull saturates with a single training mouse and remains
 228 near −0.025. The scaling effect is around 0.01 AUC, comparable to the self-versus-one-other gap
 229 of Section 4.2. Most of the resistance to drift is already present when training on a single other
 230 individual, and additional individuals refine it further. This analysis never lets the decoder see the
 231 held-out mouse, so it is unaffected by the protocol issue below.

4.3 Measuring it correctly: why the validation protocol matters

Every number reported above holds out whole recording sessions, and it is worth showing what that choice is worth. Repeating the identical analysis on the identical features with a trial-level split – shuffling trials pooled across the early sessions, so that a decoder may train on one part of a session and be tested on another – inflates the personal decoder’s within-block accuracy by +0.013 (lever pull), +0.010 (lick), +0.009 (reward) and +0.000 AUC (tone). The pooled decoder is unaffected, because it is trained on other animals and has no session in common with the test data.

The inflation therefore lands on one side of a difference between two arms (rightmost column of Table 1). For lever pull it is the entire effect: a significant-looking +0.014 ($p = 0.001$) becomes +0.003 ($p = 0.37$). For reward, which has the largest true effect, it changes little. The same reasoning applies to the count-matched control, where the natural implementation is to score the self decoder on the block it was fit on – not a cross-validated quantity at all, and enough to reverse the ordering the control exists to establish.

The mechanism behind this, and why it generalizes beyond our dataset, is set out in Section 3.3.

4.4 The effect replicates on an independent dataset

To further test whether the asymmetry is a property of that one cohort, modality and task, we repeat the analysis on the Allen Visual Behavior two-photon dataset [1]: a different laboratory, cellular two-photon rather than mesoscale widefield imaging, a visual change-detection task rather than a cued lever pull, and a 27-dimensional permutation-invariant population summary rather than a 44-parcel atlas feature. The decoded contrast is a stimulus change, so it is much less coupled to reward and licking than anything decoded on BraiDyn-BC.

The asymmetry reproduces at a comparable magnitude (Table 2). Across 23 mice the personal decoder ages by +0.024 AUC and the pooled decoder by +0.001, an asymmetry of +0.023 [0.010, 0.037], $p = 0.003$, with the mouse’s own decoder ageing more in 74% of animals. Early-block accuracy is 0.672 for the personal decoder and 0.638 for the pooled one, so the familiar ordering holds at the start; by the late block the personal decoder leads by only 0.012, about a third of its initial advantage.

The mechanism identified is not identical in both datasets; on BraiDyn-BC the gap narrows mostly because the pooled decoder *gains* over days as the animal becomes more practised (Section 4.1), while the personal decoder stays flat. Here the pooled decoder is flat (+0.001) and the entire asymmetry comes from the personal decoder losing accuracy. The shared finding within both, hence, is that the personal decoder ages faster.

The count-matched control reproduces as well. At equal training-event count the self decoder ages by +0.024, a decoder from one other mouse by +0.013 and one drawn across the other mice by +0.003. The self-minus-one gap is +0.011 [0.001, 0.022], close to the +0.010 measured on BraiDyn-BC, though on 23 animals it is marginal by the sign-flip test ($p = 0.054$); self-minus-many is +0.021 [0.010, 0.033], $p = 0.002$. The diversity ordering again holds: many other mice age less than one by -0.010 [$-0.016, -0.004$], $p = 0.003$. The result also does not depend on where the inclusion threshold is drawn, giving +0.022 at three or more sessions per mouse (24 mice), +0.023 at four (23 mice), +0.023 at five (22 mice) and +0.024 at six (18 mice).

One aspect of the analysis does not transfer between the two cohorts. Repeating the trial-level split here *lowers* the asymmetry to +0.014 ($p = 0.07$) rather than raising it. Separating the two differences on the same trials: trial shuffling still inflates the personal decoder by +0.015, as on BraiDyn-BC, but scoring one pooled AUC across sessions instead of averaging per-session values costs -0.030 ($p < 10^{-4}$) and dominates. The corrected estimate is here the *larger* one, so the replication is not an artifact of the split.

5 Discussion

A subject’s own decoder overfits idiosyncratic features of its early sessions, some of which later drift. A decoder that has never seen the subject must instead rely on components of the code that are shared across individuals and more stable over time. The count-matched control supports this reading directly: at equal data, other-individual data ages less than the subject’s own. Population structure is therefore not only a resource for cross-subject transfer but also for temporal robustness, and the

Table 2: The aging asymmetry on the original cohort and on the external validation dataset. Aging is AUC lost from early to late test with training held fixed; asymmetry is personal – pooled aging, positive when the subject’s own decoder ages more. Both rows use the same session-held-out protocol and the same count-matched control. p : paired sign-flip permutation. Per-arm aging for BraiDyn-BC is broken out by event in Table 1.

	Personal	Pooled	Asymmetry (95% CI)	p	own ages more
BraiDyn-BC, widefield, 25 mice	–	–	+0.017 [0.010, 0.024]	5×10^{-5}	73%
Allen VB, two-photon, 23 mice	+0.024	+0.001	+0.023 [0.010, 0.037]	0.003	74%

Count-matched control, aging at equal training-event count. BraiDyn-BC: self – one other mouse = +0.010; many others age less than one on all four events. Allen VB: self +0.024, one +0.013, many +0.003; self – one = +0.011, many – one = –0.010.

Best practice for longitudinal within- versus cross-subject comparisons

1. **Hold out whole recording sessions, not trials.** Only the within-subject arm can share a session between train and test, so a trial-level split inflates that arm alone. Here it inflated within-block accuracy by +0.013 (lever pull), +0.010 (lick), +0.009 (reward) and +0.000 AUC (tone), and on lever pull it accounted for the entire apparent effect: +0.014 ($p = 0.001$) became +0.003 ($p = 0.37$).
2. **Count-match the training sets before comparing arms.** A pooled decoder contains $\sim 24\times$ more training events than a personal one. Fix the training-set size to n and fit *self*, *one* and *many* on the same n , scored on the same held-out session.
3. **Never score an arm on the block it was fit on.** In the count-matched control the natural implementation scores the self decoder on its own training block – not a cross-validated quantity at all, and enough to reverse the ordering the control exists to establish.
4. **Report each arm’s aging separately, not only the difference.** The asymmetry is a difference of two quantities that can move for different reasons: on BraiDyn-BC it narrows because the pooled decoder gains, on Allen because the personal decoder loses.
5. **Average per-session scores rather than pooling one score across sessions.** On the Allen cohort, scoring a single pooled AUC across sessions instead of averaging per-session values cost –0.030 ($p < 10^{-4}$) and dominated the leakage term it was confounded with.

early-week accuracy cost of cross-subject decoding is recovered over time as the pooled decoder remains stable. Unlike within-subject recalibration, this requires no newly labelled data from the individual at test time – which is the practical point, since obtaining that data is the expensive part of deploying a decoder over months.

The effect is small in absolute terms, and we think it is worth being explicit about why we believe it anyway. It rests on 25 animals and 357 recording sessions, which is a large cohort for a longitudinal cortex-wide study, and it is supported by five checks that are not variations on one analysis. A two-block contrast and a day-by-day regression over all fifteen sessions are different estimators and select the same events. The count-matched control fixes the training-set size, so the result cannot be a data-volume artifact. Replacing the linear readout with a 1-D CNN or a GRU leaves it in place, so it is not an artifact of a weak decoder. And holding out whole sessions rather than trials removes an inflation large enough to eliminate one of the four events entirely – the effect reported here is what remains after that correction, not before it. And it reproduced in a second cohort recorded in another laboratory with a different instrument and a different task. Taken together, each of these checks could independently have removed the result, and none of them did.

Limitations. The clearest limitation is what the four events are. They are not four independent behaviors but four thresholded channels of one cued lever-pull trial, overlapping to the point where 100% of reward windows also contain licking, and differing tenfold in event count. We therefore cannot always say which behavior the effect belongs to, nor can we exclude the idea that the per-event ranking is a tracker of how much data each decoder had to overfit. A design with temporally separated, count-matched events would settle both, and is the experiment we would run next.

A second limitation is spatial scale. We measure at the mesoscale, using parcel-averaged activity, and while the two-photon replication does reach single neurons it does so only at one depth in one visual area. Whether the same asymmetry holds for single-neuron codes across cortex is therefore still open.

A third is temporal and taxonomic scope. The analysis covers a two-week interval in one species, and nothing here establishes how the asymmetry behaves over the months-to-years horizon on which a deployed decoder actually has to survive.

Finally, the trial-level inflation of Section 4.3 is worth stating as a caution rather than as a limitation of this work, because it is a limitation we do not have. It is large enough to manufacture an effect that is not there – on lever pull it manufactured the entire effect – and every number we report is measured after holding out whole sessions, not before. We raise it because any comparison of this shape is exposed to it, and the correction is the one in Box 5.

Reproducibility

BraiDyn-BC is publicly available through DANDI:001425. We release the analysis code and a streaming feature pipeline that reproduces all reported results from the parcellated traces without downloading the raw imaging data. Section 3 and the released code specify all data splits, permutation seeds, and hyperparameters. The external validation dataset of Section 4.4 is likewise public, and the released code streams it from its public bucket and reproduces that section from the same cohort selection rule.

References

- [1] Allen Institute for Brain Science. Visual Behavior — 2-photon. *Allen Brain Observatory*, 2023. <https://portal.brain-map.org/circuits-behavior/visual-behavior-2p>.
- [2] A. D. Degenhart et al. Stabilization of a brain–computer interface via the alignment of low-dimensional spaces of neural activity. *Nature Biomedical Engineering*, 4:672–685, 2020.
- [3] L. N. Driscoll et al. Dynamic reorganization of neuronal activity patterns in parietal cortex. *Cell*, 170(5):986–999, 2017.
- [4] A. Farshchian et al. Adversarial domain adaptation for stable brain–machine interfaces. *ICLR*, 2019.
- [5] B. Jarosiewicz et al. Virtual typing by people with tetraplegia using a self-calibrating intracortical brain–computer interface. *Science Translational Medicine*, 7(313):313ra179, 2015.
- [6] A. Kondo et al. BraiDyn-BC: a cortex-wide widefield calcium imaging dataset during a cued lever-pull operant task. *Scientific Data*, 12, 2025. DANDI:001425.
- [7] C. J. MacDowell and T. J. Buschman. Low-dimensional spatiotemporal dynamics underlie cortex-wide neural activity. *Current Biology*, 30(14):2665–2680, 2020.
- [8] C. J. MacDowell et al. Shared cortex-wide spatiotemporal motifs govern activity across individuals and behavioral states. *Current Biology*, 34, 2024.
- [9] S. Melbaum et al. Conserved structures of neural activity in sensorimotor cortex of freely moving rats enable cross-subject decoding. *Nature Communications*, 13:7902, 2022.
- [10] Y. Zou, K. Liu, C. Zhang, Y. Ling, F. Wang, M. Li, Y. Chen, M. Li, S. Guan, Z. He, and C. T. Li. Meta-AlignNN: a meta-learning framework for stable brain–computer interface performance across subjects, time, and tasks. *bioRxiv*, 2025. doi:10.1101/2025.04.20.649482.
- [11] S. Musall et al. Single-trial neural dynamics are dominated by richly varied movements. *Nature Neuroscience*, 22(10):1677–1686, 2019.
- [12] Y. Zhang, Y. Wang, M. Azabou, A. Andre, Z. Wang, H. Lyu, The International Brain Laboratory, E. Dyer, L. Paninski, and C. Hurwitz. Neural encoding and decoding at scale. arXiv:2504.08201, 2025.
- [13] S. M. Peterson, Z. Steine-Hanson, N. Davis, R. P. N. Rao, and B. W. Brunton. Generalized neural decoders for transfer learning across participants and recording modalities. *Journal of Neural Engineering*, 18:026014, 2021.
- [14] B. M. Karpowicz, Y. H. Ali, L. N. Wimalasena, et al. Stabilizing brain–computer interfaces through alignment of latent dynamics (NoMAD). *Nature Communications*, 16:4662, 2025.

- 352 [15] G. H. Wilson et al. Long-term unsupervised recalibration of cursor-based intracortical brain–computer
353 interfaces using a hidden Markov model. *Nature Biomedical Engineering*, 2025.
- 354 [16] M. E. Rule, T. O’Leary, and C. D. Harvey. Causes and consequences of representational drift. *Current*
355 *Opinion in Neurobiology*, 58:141–147, 2019.
- 356 [17] M. Safaie, J. C. Chang, J. Park, L. E. Miller, J. T. Dudman, M. G. Perich, and J. A. Gallego. Preserved
357 neural dynamics across animals performing similar behaviour. *Nature*, 623:765–771, 2023.
- 358 [18] M. E. Rule et al. Stable task information from an unstable neural population. *eLife*, 9:e51121, 2020.
- 359 [19] S. Schneider, J. H. Lee, and M. W. Mathis. Learnable latent embeddings for joint behavioural and neural
360 analysis (CEBRA). *Nature*, 617:360–368, 2023.
- 361 [20] C. E. Schoonover et al. Representational drift in primary olfactory cortex. *Nature*, 594:541–546, 2021.
- 362 [21] M. Hosseini, E. Erturk, S. Hashemi, and M. M. Shanechi. Cross-subject modeling for widefield calcium
363 imaging via atlas-aligned spatiotemporal tokenization. In *Proceedings of the 43rd International Conference*
364 *on Machine Learning (ICML)*, 2026. arXiv:2607.09754.